

```
In [1]: import pandas as pd
import warnings
warnings.filterwarnings('ignore')
data=pd.read_csv("/home/placement/Downloads/fiat500.csv") #reading datafile
print(data)
```

| | ID | model | engine_power | age_in_days | km | previous_owners | \ |
|------|------|--------|--------------|-------------|--------|-----------------|-----|
| 0 | 1 | lounge | 51 | 882 | 25000 | 1 | |
| 1 | 2 | pop | 51 | 1186 | 32500 | 1 | |
| 2 | 3 | sport | 74 | 4658 | 142228 | 1 | |
| 3 | 4 | lounge | 51 | 2739 | 160000 | 1 | |
| 4 | 5 | pop | 73 | 3074 | 106880 | 1 | |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 1533 | 1534 | sport | 51 | 3712 | 115280 | 1 | |
| 1534 | 1535 | lounge | 74 | 3835 | 112000 | 1 | |
| 1535 | 1536 | pop | 51 | 2223 | 60457 | 1 | |
| 1536 | 1537 | lounge | 51 | 2557 | 80750 | 1 | |
| 1537 | 1538 | pop | 51 | 1766 | 54276 | 1 | |

| | lat | lon | price |
|------|-----------|-----------|-------|
| 0 | 44.907242 | 8.611560 | 8900 |
| 1 | 45.666359 | 12.241890 | 8800 |
| 2 | 45.503300 | 11.417840 | 4200 |
| 3 | 40.633171 | 17.634609 | 6000 |
| 4 | 41.903221 | 12.495650 | 5700 |
| ... | ... | ... | ... |
| 1533 | 45.069679 | 7.704920 | 5200 |
| 1534 | 45.845692 | 8.666870 | 4600 |
| 1535 | 45.481541 | 9.413480 | 7500 |
| 1536 | 45.000702 | 7.682270 | 5990 |
| 1537 | 40.323410 | 17.568270 | 7900 |

[1538 rows x 9 columns]

```
In [2]: data.head(10)
```

```
Out[2]:
```

| | ID | model | engine_power | age_in_days | km | previous_owners | lat | lon | price |
|---|----|--------|--------------|-------------|--------|-----------------|-----------|-----------|-------|
| 0 | 1 | lounge | 51 | 882 | 25000 | 1 | 44.907242 | 8.611560 | 8900 |
| 1 | 2 | pop | 51 | 1186 | 32500 | 1 | 45.666359 | 12.241890 | 8800 |
| 2 | 3 | sport | 74 | 4658 | 142228 | 1 | 45.503300 | 11.417840 | 4200 |
| 3 | 4 | lounge | 51 | 2739 | 160000 | 1 | 40.633171 | 17.634609 | 6000 |
| 4 | 5 | pop | 73 | 3074 | 106880 | 1 | 41.903221 | 12.495650 | 5700 |
| 5 | 6 | pop | 74 | 3623 | 70225 | 1 | 45.000702 | 7.682270 | 7900 |
| 6 | 7 | lounge | 51 | 731 | 11600 | 1 | 44.907242 | 8.611560 | 10750 |
| 7 | 8 | lounge | 51 | 1521 | 49076 | 1 | 41.903221 | 12.495650 | 9190 |
| 8 | 9 | sport | 73 | 4049 | 76000 | 1 | 45.548000 | 11.549470 | 5600 |
| 9 | 10 | sport | 51 | 3653 | 89000 | 1 | 45.438301 | 10.991700 | 6000 |

```
In [3]: data.columns
```

```
Out[3]: Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
              'lat', 'lon', 'price'],  
              dtype='object')
```

```
In [4]: data.shape
```

```
Out[4]: (1538, 9)
```

In [5]: `data.describe()`

Out[5]:

| | ID | engine_power | age_in_days | km | previous_owners | lat | lon | price |
|--------------|-------------|--------------|-------------|---------------|-----------------|-------------|-------------|--------------|
| count | 1538.000000 | 1538.000000 | 1538.000000 | 1538.000000 | 1538.000000 | 1538.000000 | 1538.000000 | 1538.000000 |
| mean | 769.500000 | 51.904421 | 1650.980494 | 53396.011704 | 1.123537 | 43.541361 | 11.563428 | 8576.003901 |
| std | 444.126671 | 3.988023 | 1289.522278 | 40046.830723 | 0.416423 | 2.133518 | 2.328190 | 1939.958641 |
| min | 1.000000 | 51.000000 | 366.000000 | 1232.000000 | 1.000000 | 36.855839 | 7.245400 | 2500.000000 |
| 25% | 385.250000 | 51.000000 | 670.000000 | 20006.250000 | 1.000000 | 41.802990 | 9.505090 | 7122.500000 |
| 50% | 769.500000 | 51.000000 | 1035.000000 | 39031.000000 | 1.000000 | 44.394096 | 11.869260 | 9000.000000 |
| 75% | 1153.750000 | 51.000000 | 2616.000000 | 79667.750000 | 1.000000 | 45.467960 | 12.769040 | 10000.000000 |
| max | 1538.000000 | 77.000000 | 4658.000000 | 235000.000000 | 4.000000 | 46.795612 | 18.365520 | 11100.000000 |

Removing unwanted columns

In [6]: `data1=data.drop(columns=["ID", "lat", "lon"])`

```
In [7]: data1
```

```
Out[7]:
```

| | model | engine_power | age_in_days | km | previous_owners | price |
|------|--------|--------------|-------------|--------|-----------------|-------|
| 0 | lounge | 51 | 882 | 25000 | 1 | 8900 |
| 1 | pop | 51 | 1186 | 32500 | 1 | 8800 |
| 2 | sport | 74 | 4658 | 142228 | 1 | 4200 |
| 3 | lounge | 51 | 2739 | 160000 | 1 | 6000 |
| 4 | pop | 73 | 3074 | 106880 | 1 | 5700 |
| ... | ... | ... | ... | ... | ... | ... |
| 1533 | sport | 51 | 3712 | 115280 | 1 | 5200 |
| 1534 | lounge | 74 | 3835 | 112000 | 1 | 4600 |
| 1535 | pop | 51 | 2223 | 60457 | 1 | 7500 |
| 1536 | lounge | 51 | 2557 | 80750 | 1 | 5990 |
| 1537 | pop | 51 | 1766 | 54276 | 1 | 7900 |

1538 rows × 6 columns

```
In [8]: data1=pd.get_dummies(data1)
```

```
In [9]: data1
```

```
Out[9]:
```

| | engine_power | age_in_days | km | previous_owners | price | model_lounge | model_pop | model_sport |
|------|--------------|-------------|--------|-----------------|-------|--------------|-----------|-------------|
| 0 | 51 | 882 | 25000 | 1 | 8900 | 1 | 0 | 0 |
| 1 | 51 | 1186 | 32500 | 1 | 8800 | 0 | 1 | 0 |
| 2 | 74 | 4658 | 142228 | 1 | 4200 | 0 | 0 | 1 |
| 3 | 51 | 2739 | 160000 | 1 | 6000 | 1 | 0 | 0 |
| 4 | 73 | 3074 | 106880 | 1 | 5700 | 0 | 1 | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1533 | 51 | 3712 | 115280 | 1 | 5200 | 0 | 0 | 1 |
| 1534 | 74 | 3835 | 112000 | 1 | 4600 | 1 | 0 | 0 |
| 1535 | 51 | 2223 | 60457 | 1 | 7500 | 0 | 1 | 0 |
| 1536 | 51 | 2557 | 80750 | 1 | 5990 | 1 | 0 | 0 |
| 1537 | 51 | 1766 | 54276 | 1 | 7900 | 0 | 1 | 0 |

1538 rows × 8 columns

```
In [10]: data1.shape
```

```
Out[10]: (1538, 8)
```

remove the actual value from the dataframe

```
In [11]: y=data1['price']  
x=data1.drop(columns='price')
```

In [12]:

x

Out[12]:

| | engine_power | age_in_days | km | previous_owners | model_lounge | model_pop | model_sport |
|------|--------------|-------------|--------|-----------------|--------------|-----------|-------------|
| 0 | 51 | 882 | 25000 | 1 | 1 | 0 | 0 |
| 1 | 51 | 1186 | 32500 | 1 | 0 | 1 | 0 |
| 2 | 74 | 4658 | 142228 | 1 | 0 | 0 | 1 |
| 3 | 51 | 2739 | 160000 | 1 | 1 | 0 | 0 |
| 4 | 73 | 3074 | 106880 | 1 | 0 | 1 | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 1533 | 51 | 3712 | 115280 | 1 | 0 | 0 | 1 |
| 1534 | 74 | 3835 | 112000 | 1 | 1 | 0 | 0 |
| 1535 | 51 | 2223 | 60457 | 1 | 0 | 1 | 0 |
| 1536 | 51 | 2557 | 80750 | 1 | 1 | 0 | 0 |
| 1537 | 51 | 1766 | 54276 | 1 | 0 | 1 | 0 |

1538 rows × 7 columns

In [13]:

y

Out[13]:

```

0      8900
1      8800
2      4200
3      6000
4      5700
...
1533   5200
1534   4600
1535   7500
1536   5990
1537   7900

```

Name: price, Length: 1538, dtype: int64

```
In [14]: #!pip install scikit-learn
```

split the data into training set and testing set

```
In [15]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)# 66% and 33%
```

```
In [16]: x_test
```

```
Out[16]:
```

| | engine_power | age_in_days | km | previous_owners | model_lounge | model_pop | model_sport |
|------|--------------|-------------|--------|-----------------|--------------|-----------|-------------|
| 481 | 51 | 3197 | 120000 | 2 | 0 | 1 | 0 |
| 76 | 62 | 2101 | 103000 | 1 | 0 | 1 | 0 |
| 1502 | 51 | 670 | 32473 | 1 | 1 | 0 | 0 |
| 669 | 51 | 913 | 29000 | 1 | 1 | 0 | 0 |
| 1409 | 51 | 762 | 18800 | 1 | 1 | 0 | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 291 | 51 | 701 | 22000 | 1 | 1 | 0 | 0 |
| 596 | 51 | 3347 | 85500 | 1 | 0 | 1 | 0 |
| 1489 | 51 | 366 | 22148 | 1 | 0 | 1 | 0 |
| 1436 | 51 | 1797 | 61000 | 1 | 1 | 0 | 0 |
| 575 | 51 | 366 | 19112 | 1 | 1 | 0 | 0 |

508 rows × 7 columns

In [17]: x_train

Out[17]:

| | engine_power | age_in_days | km | previous_owners | model_lounge | model_pop | model_sport |
|------|--------------|-------------|--------|-----------------|--------------|-----------|-------------|
| 527 | 51 | 425 | 13111 | 1 | 1 | 0 | 0 |
| 129 | 51 | 1127 | 21400 | 1 | 1 | 0 | 0 |
| 602 | 51 | 2039 | 57039 | 1 | 0 | 1 | 0 |
| 331 | 51 | 1155 | 40700 | 1 | 1 | 0 | 0 |
| 323 | 51 | 425 | 16783 | 1 | 1 | 0 | 0 |
| ... | ... | ... | ... | ... | ... | ... | ... |
| 1130 | 51 | 1127 | 24000 | 1 | 1 | 0 | 0 |
| 1294 | 51 | 852 | 30000 | 1 | 1 | 0 | 0 |
| 860 | 51 | 3409 | 118000 | 1 | 0 | 1 | 0 |
| 1459 | 51 | 762 | 16700 | 1 | 1 | 0 | 0 |
| 1126 | 51 | 701 | 39207 | 1 | 1 | 0 | 0 |

1030 rows × 7 columns

In [18]: y_test.head(5)

Out[18]:

| | |
|------|------|
| 481 | 7900 |
| 76 | 7900 |
| 1502 | 9400 |
| 669 | 8500 |
| 1409 | 9700 |

Name: price, dtype: int64


```
In [19]: y_train.head(5)
```

```
Out[19]: 527    9990
         129    9500
         602    7590
         331    8750
         323    9100
         Name: price, dtype: int64
```

LinearRegression

```
In [20]: from sklearn.linear_model import LinearRegression
         reg=LinearRegression()#creating object of LinearRegression
         reg.fit(x_train,y_train)#training and fitting
```

```
Out[20]: LinearRegression()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [55]: y_pred_regression=reg.predict(x_test) #predict the price using x_test data
```

In [56]: `y_pred_regression`

Out[56]: `array([5867.6503378 , 7133.70142341, 9866.35776216, 9723.28874535,`
`10039.59101162, 9654.07582608, 9673.14563045, 10118.70728123,`
`9903.85952664, 9351.55828437, 10434.34963575, 7732.26255693,`
`7698.67240131, 6565.95240435, 9662.90103518, 10373.20344286,`
`9599.94844451, 7699.34400418, 4941.33017994, 10455.2719478 ,`
`10370.51555682, 10391.60424404, 7529.06622456, 9952.37340054,`
`7006.13845729, 9000.1780961 , 4798.36770637, 6953.10376491,`
`7810.39767825, 9623.80497535, 7333.52158317, 5229.18705519,`
`5398.21541073, 5157.65652129, 8948.63632836, 5666.62365159,`
`9822.1231461 , 8258.46551788, 6279.2040404 , 8457.38443276,`
`9773.86444066, 6767.04074749, 9182.99904787, 10210.05195479,`
`8694.90545226, 10328.43369248, 9069.05761443, 8866.7826029 ,`
`7058.39787506, 9073.33877162, 9412.68162121, 10293.69451263,`
`10072.49011135, 6748.5794244 , 9785.95841801, 9354.09969973,`
`9507.9444386 , 10443.01608254, 9795.31884316, 7197.84932877,`
`10108.31707235, 7009.6597206 , 9853.90699412, 7146.87414965,`
`6417.69133992, 9996.97382441, 9781.18795953, 8515.83255277,`
`8456.30006203, 6499.76668237, 7768.57829985, 6832.86406122,`
`8347.96113362, 10439.02404036, 7356.43463051, 8562.56562053,`
`8820.78555100, 10025.82571520, 7270.77100000, 8411.45004000`

In [57]: `from sklearn.metrics import r2_score #to know the efficiency bw the predicted price`
`r2_score(y_test,y_pred_regression)`

Out[57]: `0.8415526986865394`

In [58]: `from sklearn.metrics import mean_squared_error#calaculating mse`
`mean_squared_error(y_test,y_pred_regression)`

Out[58]: `581887.727391353`

In [59]: `import math`
`a=581887.727391353`
`print(math.sqrt(a))`

`762.8156575420782`

```
In [60]: y_test.head(10)
```

```
Out[60]: 481      7900
         76      7900
         1502    9400
         669    8500
         1409    9700
         1414    9900
         1089    9900
         1507    9950
         970   10700
         1198    8999
         Name: price, dtype: int64
```

```
In [61]: results=pd.DataFrame(columns=['Price','Predicted']) #create datafame for price and predicted
         results['Price']=y_test
         results['Predicted']=y_pred_regression
         results=results.reset_index() #remove the index as ID values
         results['id']=results.index
```

In [62]: results

Out[62]:

| | index | Price | Predicted | id |
|-----|-------|-------|--------------|-----|
| 0 | 481 | 7900 | 5867.650338 | 0 |
| 1 | 76 | 7900 | 7133.701423 | 1 |
| 2 | 1502 | 9400 | 9866.357762 | 2 |
| 3 | 669 | 8500 | 9723.288745 | 3 |
| 4 | 1409 | 9700 | 10039.591012 | 4 |
| ... | ... | ... | ... | ... |
| 503 | 291 | 10900 | 10032.665135 | 503 |
| 504 | 596 | 5699 | 6281.536277 | 504 |
| 505 | 1489 | 9500 | 9986.327508 | 505 |
| 506 | 1436 | 6990 | 8381.517020 | 506 |
| 507 | 575 | 10900 | 10371.142553 | 507 |

508 rows × 4 columns

In [63]: results["Difference"]=results.apply(lambda x:x.Price-x.Predicted,axis=1)#add the column for difference b/w t

In [64]:

```
results.head(10)
```

Out[64]:

| | index | Price | Predicted | id | Difference |
|---|-------|-------|--------------|----|--------------|
| 0 | 481 | 7900 | 5867.650338 | 0 | 2032.349662 |
| 1 | 76 | 7900 | 7133.701423 | 1 | 766.298577 |
| 2 | 1502 | 9400 | 9866.357762 | 2 | -466.357762 |
| 3 | 669 | 8500 | 9723.288745 | 3 | -1223.288745 |
| 4 | 1409 | 9700 | 10039.591012 | 4 | -339.591012 |
| 5 | 1414 | 9900 | 9654.075826 | 5 | 245.924174 |
| 6 | 1089 | 9900 | 9673.145630 | 6 | 226.854370 |
| 7 | 1507 | 9950 | 10118.707281 | 7 | -168.707281 |
| 8 | 970 | 10700 | 9903.859527 | 8 | 796.140473 |
| 9 | 1198 | 8999 | 9351.558284 | 9 | -352.558284 |

In [65]: results


Out[65]:

| | index | Price | Predicted | id | Difference |
|-----|-------|-------|--------------|-----|--------------|
| 0 | 481 | 7900 | 5867.650338 | 0 | 2032.349662 |
| 1 | 76 | 7900 | 7133.701423 | 1 | 766.298577 |
| 2 | 1502 | 9400 | 9866.357762 | 2 | -466.357762 |
| 3 | 669 | 8500 | 9723.288745 | 3 | -1223.288745 |
| 4 | 1409 | 9700 | 10039.591012 | 4 | -339.591012 |
| ... | ... | ... | ... | ... | ... |
| 503 | 291 | 10900 | 10032.665135 | 503 | 867.334865 |
| 504 | 596 | 5699 | 6281.536277 | 504 | -582.536277 |
| 505 | 1489 | 9500 | 9986.327508 | 505 | -486.327508 |
| 506 | 1436 | 6990 | 8381.517020 | 506 | -1391.517020 |
| 507 | 575 | 10900 | 10371.142553 | 507 | 528.857447 |

508 rows × 5 columns

plot the data using seaborn and matplotlib libraries

```
In [33]: import seaborn as sns
import matplotlib.pyplot as plt
sns.lineplot(x='id',y='Price',data=results.head(50))      #actual color=blue
sns.lineplot(x='id',y='Predicted',data=results.head(50)) #predicted color=orange
plt.show()
```



Ridge_Regression

```
In [34]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import Ridge
alpha=[1e-15,1e-10,1e-4,1e-3,1e-2,15,10,20,30]
ridge=Ridge()                                #creating an object for Ridge
parameters={'alpha':alpha}
ridge_regressor=GridSearchCV(ridge,parameters)
ridge_regressor.fit(x_train,y_train)         #training and fitting
```

```
Out[34]: GridSearchCV(estimator=Ridge(),
                      param_grid={'alpha': [1e-15, 1e-10, 0.0001, 0.001, 0.01, 15, 10,
                                             20, 30]})
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [35]: ridge_regressor.best_params_
```

```
Out[35]: {'alpha': 30}
```

```
In [66]: ridge=Ridge(alpha=30)
ridge.fit(x_train,y_train)
y_pred_ridge=ridge.predict(x_test)
```

```
In [67]: from sklearn.metrics import r2_score #to know the efficiency of the predicted price
r2_score(y_test,y_pred_ridge)
```

```
Out[67]: 0.8421969385523054
```



```
In [68]: from sklearn.metrics import mean_squared_error
Ridge_Error=mean_squared_error(y_pred_ridge,y_test)
Ridge_Error
```

Out[68]: 579521.7970897449

```
In [69]: results=pd.DataFrame(columns=['Actual','Predicted'])    #create the dataframe for actual and predicted values
results['Actual']=y_test
results['Predicted']=y_pred_ridge
results=results.reset_index()    #remove the index as ID values
results['id']=results.index
```

```
In [70]: results
```

Out[70]:

| | index | Actual | Predicted | id |
|-----|-------|--------|--------------|-----|
| 0 | 481 | 7900 | 5869.741155 | 0 |
| 1 | 76 | 7900 | 7149.563327 | 1 |
| 2 | 1502 | 9400 | 9862.785355 | 2 |
| 3 | 669 | 8500 | 9719.283532 | 3 |
| 4 | 1409 | 9700 | 10035.895686 | 4 |
| ... | ... | ... | ... | ... |
| 503 | 291 | 10900 | 10029.070743 | 503 |
| 504 | 596 | 5699 | 6297.833772 | 504 |
| 505 | 1489 | 9500 | 10008.285472 | 505 |
| 506 | 1436 | 6990 | 8375.789449 | 506 |
| 507 | 575 | 10900 | 10368.170257 | 507 |

508 rows × 4 columns

```
In [71]: results["Difference"]=results.apply(lambda x:x.Actual-x.Predicted,axis=1)#add the column for difference b/w
```


In [72]: results

Out[72]:

| | index | Actual | Predicted | id | Difference |
|-----|-------|--------|--------------|-----|--------------|
| 0 | 481 | 7900 | 5869.741155 | 0 | 2030.258845 |
| 1 | 76 | 7900 | 7149.563327 | 1 | 750.436673 |
| 2 | 1502 | 9400 | 9862.785355 | 2 | -462.785355 |
| 3 | 669 | 8500 | 9719.283532 | 3 | -1219.283532 |
| 4 | 1409 | 9700 | 10035.895686 | 4 | -335.895686 |
| ... | ... | ... | ... | ... | ... |
| 503 | 291 | 10900 | 10029.070743 | 503 | 870.929257 |
| 504 | 596 | 5699 | 6297.833772 | 504 | -598.833772 |
| 505 | 1489 | 9500 | 10008.285472 | 505 | -508.285472 |
| 506 | 1436 | 6990 | 8375.789449 | 506 | -1385.789449 |
| 507 | 575 | 10900 | 10368.170257 | 507 | 531.829743 |

508 rows × 5 columns

```
In [73]: import seaborn as sns
import matplotlib.pyplot as plt
sns.lineplot(x='id',y='Actual',data=results.head(50))
sns.lineplot(x='id',y='Predicted',data=results.head(50))
plt.show()
```



ElastiNet model

```
In [44]: from sklearn.linear_model import ElasticNet
from sklearn.model_selection import GridSearchCV
elastic=ElasticNet()                                #creating an object for ElasticNet
parameters={'alpha':[1e-15,1e-10,1e-8,1e-4,1e-3,1e-2,15,10,20]}
elastic_regressor=GridSearchCV(elastic,parameters)
elastic_regressor.fit(x_train,y_train)              #training and fitting
```

```
Out[44]: GridSearchCV(estimator=ElasticNet(),
                      param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.0001, 0.001, 0.01, 15,
                                             10, 20]})
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [45]: elastic_regressor.best_params_
```

```
Out[45]: {'alpha': 0.01}
```

```
In [74]: elastic=ElasticNet(alpha=0.01)
elastic.fit(x_train,y_train)
y_pred_enet=elastic.predict(x_test)
```

```
In [75]: from sklearn.metrics import mean_squared_error
elastic_Error=mean_squared_error(y_pred_enet,y_test)
elastic_Error
```

```
Out[75]: 581390.7642825295
```

```
In [76]: from sklearn.metrics import r2_score #to know the efficiency of the predicted price
r2_score(y_test,y_pred_enet)
```

```
Out[76]: 0.841688021120299
```

```
In [77]: results=pd.DataFrame(columns=['Actual','Predicted']) #create the dataframe for actual and predicted values
results['Actual']=y_test
results['Predicted']=y_pred_enet
results=results.reset_index() #remove the index as ID values
results['id']=results.index
```

```
In [78]: results
```

```
Out[78]:
```

| | index | Actual | Predicted | id |
|-----|-------|--------|--------------|-----|
| 0 | 481 | 7900 | 5867.742075 | 0 |
| 1 | 76 | 7900 | 7136.527402 | 1 |
| 2 | 1502 | 9400 | 9865.726723 | 2 |
| 3 | 669 | 8500 | 9722.573593 | 3 |
| 4 | 1409 | 9700 | 10038.936496 | 4 |
| ... | ... | ... | ... | ... |
| 503 | 291 | 10900 | 10032.030157 | 503 |
| 504 | 596 | 5699 | 6284.484674 | 504 |
| 505 | 1489 | 9500 | 9990.379510 | 505 |
| 506 | 1436 | 6990 | 8380.465651 | 506 |
| 507 | 575 | 10900 | 10370.628731 | 507 |

508 rows × 4 columns

```
In [79]: results["Difference"]=results['Actual']-results['Predicted']#add the column for difference b/w the actual and predicted
```

```
In [80]: results
```


```
Out[80]:
```

| | index | Actual | Predicted | id | Difference |
|-----|-------|--------|--------------|-----|--------------|
| 0 | 481 | 7900 | 5867.742075 | 0 | 2032.257925 |
| 1 | 76 | 7900 | 7136.527402 | 1 | 763.472598 |
| 2 | 1502 | 9400 | 9865.726723 | 2 | -465.726723 |
| 3 | 669 | 8500 | 9722.573593 | 3 | -1222.573593 |
| 4 | 1409 | 9700 | 10038.936496 | 4 | -338.936496 |
| ... | ... | ... | ... | ... | ... |
| 503 | 291 | 10900 | 10032.030157 | 503 | 867.969843 |
| 504 | 596 | 5699 | 6284.484674 | 504 | -585.484674 |
| 505 | 1489 | 9500 | 9990.379510 | 505 | -490.379510 |
| 506 | 1436 | 6990 | 8380.465651 | 506 | -1390.465651 |
| 507 | 575 | 10900 | 10370.628731 | 507 | 529.371269 |

508 rows × 5 columns

Plot the data using seaborn and matplotlib libraries

```
In [54]: import seaborn as sns
import matplotlib.pyplot as plt
sns.lineplot(x='id',y='Actual',data=results.head(50))
sns.lineplot(x='id',y='Predicted',data=results.head(50))
plt.show()
```



In []:

In []:

In []:

In []: