

```
In [1]: import pandas as pd
import warnings
warnings.filterwarnings('ignore')
data=pd.read_csv("/home/placement/Downloads/fiat500.csv") #reading datafile
print(data)
```

	ID	model	engine_power	age_in_days	km	previous_owners	\
0	1	lounge	51	882	25000	1	
1	2	pop	51	1186	32500	1	
2	3	sport	74	4658	142228	1	
3	4	lounge	51	2739	160000	1	
4	5	pop	73	3074	106880	1	
...	...	...	...	...	...	...	...
1533	1534	sport	51	3712	115280	1	
1534	1535	lounge	74	3835	112000	1	
1535	1536	pop	51	2223	60457	1	
1536	1537	lounge	51	2557	80750	1	
1537	1538	pop	51	1766	54276	1	

	lat	lon	price
0	44.907242	8.611560	8900
1	45.666359	12.241890	8800
2	45.503300	11.417840	4200
3	40.633171	17.634609	6000
4	41.903221	12.495650	5700
...	...	...	...
1533	45.069679	7.704920	5200
1534	45.845692	8.666870	4600
1535	45.481541	9.413480	7500
1536	45.000702	7.682270	5990
1537	40.323410	17.568270	7900

[1538 rows x 9 columns]

```
In [2]: data.head(10)
```

```
Out[2]:
```

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700
5	6	pop	74	3623	70225	1	45.000702	7.682270	7900
6	7	lounge	51	731	11600	1	44.907242	8.611560	10750
7	8	lounge	51	1521	49076	1	41.903221	12.495650	9190
8	9	sport	73	4049	76000	1	45.548000	11.549470	5600
9	10	sport	51	3653	89000	1	45.438301	10.991700	6000

```
In [3]: data.columns
```

```
Out[3]: Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
              'lat', 'lon', 'price'],  
              dtype='object')
```

```
In [4]: data.shape
```

```
Out[4]: (1538, 9)
```

In [5]: `data.describe()`

Out[5]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
<b>count</b>	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
<b>mean</b>	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.563428	8576.003901
<b>std</b>	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.328190	1939.958641
<b>min</b>	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.245400	2500.000000
<b>25%</b>	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.505090	7122.500000
<b>50%</b>	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.869260	9000.000000
<b>75%</b>	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.769040	10000.000000
<b>max</b>	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.365520	11100.000000

## Removing unwanted columns

In [6]: `data1=data.drop(columns=["ID", "lat", "lon"])`

```
In [7]: data1
```

```
Out[7]:
```

	model	engine_power	age_in_days	km	previous_owners	price
0	lounge	51	882	25000	1	8900
1	pop	51	1186	32500	1	8800
2	sport	74	4658	142228	1	4200
3	lounge	51	2739	160000	1	6000
4	pop	73	3074	106880	1	5700
...	...	...	...	...	...	...
1533	sport	51	3712	115280	1	5200
1534	lounge	74	3835	112000	1	4600
1535	pop	51	2223	60457	1	7500
1536	lounge	51	2557	80750	1	5990
1537	pop	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [8]: data1=pd.get_dummies(data1)
```

```
In [9]: data1
```

```
Out[9]:
```

	engine_power	age_in_days	km	previous_owners	price	model_lounge	model_pop	model_sport
0	51	882	25000	1	8900	1	0	0
1	51	1186	32500	1	8800	0	1	0
2	74	4658	142228	1	4200	0	0	1
3	51	2739	160000	1	6000	1	0	0
4	73	3074	106880	1	5700	0	1	0
...	...	...	...	...	...	...	...	...
1533	51	3712	115280	1	5200	0	0	1
1534	74	3835	112000	1	4600	1	0	0
1535	51	2223	60457	1	7500	0	1	0
1536	51	2557	80750	1	5990	1	0	0
1537	51	1766	54276	1	7900	0	1	0

1538 rows × 8 columns

## remove the actual value from the dataframe

```
In [10]: y=data1['price']  
x=data1.drop(columns='price')
```

In [11]:

x

Out[11]:

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
0	51	882	25000	1	1	0	0
1	51	1186	32500	1	0	1	0
2	74	4658	142228	1	0	0	1
3	51	2739	160000	1	1	0	0
4	73	3074	106880	1	0	1	0
...	...	...	...	...	...	...	...
1533	51	3712	115280	1	0	0	1
1534	74	3835	112000	1	1	0	0
1535	51	2223	60457	1	0	1	0
1536	51	2557	80750	1	1	0	0
1537	51	1766	54276	1	0	1	0

1538 rows × 7 columns

In [12]:

y

Out[12]:

0	8900
1	8800
2	4200
3	6000
4	5700
...	...
1533	5200
1534	4600
1535	7500
1536	5990
1537	7900

Name: price, Length: 1538, dtype: int64

## split the data into training set and testing set

```
In [13]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)# 66% and 33%
```

```
In [14]: x_test
```

Out[14]:

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
481	51	3197	120000	2	0	1	0
76	62	2101	103000	1	0	1	0
1502	51	670	32473	1	1	0	0
669	51	913	29000	1	1	0	0
1409	51	762	18800	1	1	0	0
...	...	...	...	...	...	...	...
291	51	701	22000	1	1	0	0
596	51	3347	85500	1	0	1	0
1489	51	366	22148	1	0	1	0
1436	51	1797	61000	1	1	0	0
575	51	366	19112	1	1	0	0

508 rows × 7 columns

```
In [15]: x_train
```

```
Out[15]:
```

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
527	51	425	13111	1	1	0	0
129	51	1127	21400	1	1	0	0
602	51	2039	57039	1	0	1	0
331	51	1155	40700	1	1	0	0
323	51	425	16783	1	1	0	0
...	...	...	...	...	...	...	...
1130	51	1127	24000	1	1	0	0
1294	51	852	30000	1	1	0	0
860	51	3409	118000	1	0	1	0
1459	51	762	16700	1	1	0	0
1126	51	701	39207	1	1	0	0

1030 rows × 7 columns

```
In [16]: y_test.head(5)
```

```
Out[16]: 481    7900
76      7900
1502    9400
669     8500
1409    9700
Name: price, dtype: int64
```



```
In [17]: y_train.head(5)
```

```
Out[17]: 527    9990
         129    9500
         602    7590
         331    8750
         323    9100
         Name: price, dtype: int64
```

## LinearRegression

```
In [18]: from sklearn.linear_model import LinearRegression
         reg=LinearRegression()#creating object of LinearRegression
         reg.fit(x_train,y_train)#training and fitting
```

```
Out[18]: ▼ LinearRegression
         LinearRegression()
```

```
In [19]: y_pred=reg.predict(x_test) #predict the price using x_test data
```

In [20]: y\_pred

Out[20]: array([ 5867.6503378 , 7133.70142341, 9866.35776216, 9723.28874535,  
10039.59101162, 9654.07582608, 9673.14563045, 10118.70728123,  
9903.85952664, 9351.55828437, 10434.34963575, 7732.26255693,  
7698.67240131, 6565.95240435, 9662.90103518, 10373.20344286,  
9599.94844451, 7699.34400418, 4941.33017994, 10455.2719478 ,  
10370.51555682, 10391.60424404, 7529.06622456, 9952.37340054,  
7006.13845729, 9000.1780961 , 4798.36770637, 6953.10376491,  
7810.39767825, 9623.80497535, 7333.52158317, 5229.18705519,  
5398.21541073, 5157.65652129, 8948.63632836, 5666.62365159,  
9822.1231461 , 8258.46551788, 6279.2040404 , 8457.38443276,  
9773.86444066, 6767.04074749, 9182.99904787, 10210.05195479,  
8694.90545226, 10328.43369248, 9069.05761443, 8866.7826029 ,  
7058.39787506, 9073.33877162, 9412.68162121, 10293.69451263,  
10072.49011135, 6748.5794244 , 9785.95841801, 9354.09969973,  
9507.9444386 , 10443.01608254, 9795.31884316, 7197.84932877,  
10108.31707235, 7009.6597206 , 9853.90699412, 7146.87414965,  
6417.69133992, 9996.97382441, 9781.18795953, 8515.83255277,  
8456.30006203, 6499.76668237, 7768.57829985, 6832.86406122,  
8347.96113362, 10439.02404036, 7356.43463051, 8562.56562053,  
8820.78555100, 10025.82571520, 7270.77100022, 8411.45004000])

In [21]: `from sklearn.metrics import r2_score #to know the efficiency bw the predicted price`  
`r2_score(y_test,y_pred)`

Out[21]: 0.8415526986865394

In [22]: `from sklearn.metrics import mean_squared_error#calaculating mse`  
`mean_squared_error(y_test,y_pred)`

Out[22]: 581887.727391353

```
In [24]: y_test.head(10)
```

```
Out[24]: 481      7900
         76      7900
         1502    9400
         669    8500
         1409    9700
         1414    9900
         1089    9900
         1507    9950
         970   10700
         1198    8999
         Name: price, dtype: int64
```

```
In [25]: results=pd.DataFrame(columns=['Price', 'Predicted']) #create dataframe for price and predicted
         results['Price']=y_test
         results['Predicted']=y_pred
         results=results.reset_index() #remove the index as ID values
         results['id']=results.index
```

In [26]: results

Out[26]:

	index	Price	Predicted	id
0	481	7900	5867.650338	0
1	76	7900	7133.701423	1
2	1502	9400	9866.357762	2
3	669	8500	9723.288745	3
4	1409	9700	10039.591012	4
...	...	...	...	...
503	291	10900	10032.665135	503
504	596	5699	6281.536277	504
505	1489	9500	9986.327508	505
506	1436	6990	8381.517020	506
507	575	10900	10371.142553	507

508 rows × 4 columns

In [27]: results["Difference"]=results.apply(**lambda** x:x.Price-x.Predicted,axis=1)*#add the column for difference b/w t*

In [28]: results

Out[28]:

	index	Price	Predicted	id	Difference
0	481	7900	5867.650338	0	2032.349662
1	76	7900	7133.701423	1	766.298577
2	1502	9400	9866.357762	2	-466.357762
3	669	8500	9723.288745	3	-1223.288745
4	1409	9700	10039.591012	4	-339.591012
...	...	...	...	...	...
503	291	10900	10032.665135	503	867.334865
504	596	5699	6281.536277	504	-582.536277
505	1489	9500	9986.327508	505	-486.327508
506	1436	6990	8381.517020	506	-1391.517020
507	575	10900	10371.142553	507	528.857447

508 rows × 5 columns

## plot the data using seaborn and matplotlib libraries

```
In [29]: import seaborn as sns
import matplotlib.pyplot as plt
sns.lineplot(x='id',y='Price',data=results.head(50))      #actual color=blue
sns.lineplot(x='id',y='Predicted',data=results.head(50)) #predicted color=orange
plt.show()
```



